

Deadlocks

Module 21

Resources

- Computers are really just a collection of resources...
 - The Operating System allocates / grants access to resources
 - Resources might be hardware
 - Resources might be software (the file system, a mutex, etc.)

Resource Acquisition

- Through system calls:
 - Request the resource
 - Use the resource
 - Release the resource

```
semaphore r1;
```

```
// work with resource
```

```
down(&r1);  
use_resource();  
up(&r2);
```

```
file_descriptor fd;
```

```
// use device
```

```
fd = open(device_num);  
use_device();  
close(fd);
```

Deadlocks

- A deadlock occurs when:
 - Each process in a set is waiting for an event
 - Those event can only be caused by another process in the set.
- Deadlocks occur within a set of process, and at least two shared resources

Deadlock - basic example

Process A:

Give me Resource 1

Give me Resource 2

Work on resource 1 & 2

Process A:

Give me Resource 2

Give me Resource 1

Work on resource 1 & 2

- If Process A gets Resource 1, and is then interrupted, Process B may acquire Resource 2
- Neither can run... ever again.


Requirements

A set of process, and a set of policies, must meet the following conditions:

- 1) Mutual Exclusion is Guaranteed
- 2) Hold & Wait is Permitted
- 3) No Preemption of Resources is allowed
- 4) Circular Wait is Permitted

Modeling Resource Allocation

- Deadlocks can be defined as *cycles* in a graph
 - Resources & Processes are Nodes
 - Assignment Edges
 - Request Edges



Computers are good at finding cycles in graphs!

OS Responsibility

What can an OS do?

Option 1: Pretend this will never happen.

Most general purpose OS's actually do this



Handling Deadlocks

- In some systems (long running, critical uptime), we may need to handle deadlocks
- Often individual programs handle deadlock within their own threads/resources
- Eliminate them:
 - Prevention algorithms (attack requirements as matter of policy)
 - Avoid them (careful planning)
- Detect and Recover

Attacking Requirements

- Do not allow mutual exclusion
- or ● Do not allow a process to ask for a resource while holding another
- or ● Allow resources to be taken away at any time
- or ● Detect circular wait and fail on resource request

Avoidance

- We can model “hypothetical” situations:
 - If I grant access to this resource, will there be a deadlock?
 - If so, cause the “open” call to fail
- We can hold “reserve” resources to avoid unsafe situations:
 - Banker’s algorithm - don’t allocate if there is no way of recovering from “worst case” - which is that everyone opens resources

Banker's Algorithm

12 Tape Drives, 2 currently free

Process	Max resource possible to request	Current resources acquired
P0	10	5
P1	4	2
P2	9	3

The system is “safe” - if each process asks for ALL resources it could ask for, the situation can resolve itself (P1)

Banker's Algorithm


12 Tape Drives, 1 currently free

Process	Max resource possible to request	Current resources acquired
P0	10	5
P1	4	2
P2	9	4

If P2 asks for (and gets) 1 more drive, it would put us in an **unsafe** state
Disallow P2 from acquiring the resource

Banker's Algorithm

- It is a bit “conservative”
 - *Just like banks, most OS's throw caution to the wind and charge towards the unsafe assuming they won't be around to clean up the mess...*
- Algorithm is flawed - requires us to know how many total resources a process needs, in advance!



Recovery

- Avoiding Deadlocks usually put onerous restrictions (too conservative) on system
- How about we let them happen, and then recover?
 - Terminate offending process?
 - Preempt the resource?

Why do we care?

- If most OS's don't deal with deadlocks... why are we studying them?
- Deadlocks most often occur between threads ***within the same process.***
- As a user programmer, you probably do know how many resources each thread may need, and you probably can accurately implement avoidance algorithms!

